

CO2 Emission by Vehicles

September 17, 2024

```
<h1 style='padding: 20px;
      color:white;
      text-align:center;'>
    PREDICTING CO2 EMISSIONS WITH ML LINEAR MODELS
</h1>
</div>
```

Linear Regression

Amarja Kumbhar



0.1 About the Project

Objective: The goal of this project is perform an EDA to prepare the data for a Machine Learning Model, focusing on building a Linear Model.

The project consists of two main phases:

1. Exploratory Data Analysis (EDA):

- For the dataset of Co2 Emission by Vehicle in Canada, performed an Exploratory Data Analysis (EDA) to understand the general structure of the dataset, summarize key

statistical insights, and explore relationships between independent variables and the target variable.

2. **Machine Learning Model:** Preparing the data and building models, including Simple and Multiple Linear Regression, Polynomial Regression, and Regularization techniques.

- **Simple Linear Regression Model:**
 - **Multiple Linear Regression Model:**
 - **Polynomial Regression Model:**
 - **Scaling the Data:**
 - **Final Model and Prediction:** The final model was built and predictions were made.
- This dataset contains information about various vehicles' carbon dioxide (CO2) emissions and fuel consumption.
 - In the context of Machine Learning (ML), this dataset is often used to predict CO2 emissions based on vehicle characteristics or to analyze fuel efficiency of vehicles.
 - The goal could be to predict CO2 emissions or fuel consumption based on the features of the vehicles.
 - There are total 7385 rows and 12 columns.
 - This dataset captures the details of how CO2 emissions by a vehicle can vary with the different features. The dataset has been taken from Canada Government official open data website. This is a compiled version. This contains data over a period of 7 years.

Data Source: <https://www.kaggle.com/datasets/debajyotipodder/co2-emission-by-vehicles>

The columns in the dataset can be described as follows:

1. **Make:** The brand of the vehicle.
2. **Model:** The model of the vehicle.
3. **Vehicle Class:** The class of the vehicle (e.g., compact, SUV).
4. **Engine Size(L):** The engine size in liters.
5. **Cylinders:** The number of cylinders in the engine.
6. **Transmission:** The type of transmission (e.g., automatic, manual).
7. **Fuel Type:** The type of fuel used (e.g., gasoline, diesel).
8. **Fuel Consumption City (L/100 km):** Fuel consumption in the city (liters per 100 kilometers).
9. **Fuel Consumption Hwy (L/100 km):** Highway (out-of-city) fuel consumption.
10. **Fuel Consumption Comb (L/100 km):** Combined (city and highway) fuel consumption.
11. **Fuel Consumption Comb (mpg):** Combined fuel consumption in miles per gallon. (efficiency-> less fuel long way)
12. **CO2 Emissions(g/km):** CO2 emissions in grams per kilometer.

NOTE:

11. **Fuel Consumption Comb (mpg):**

- High mpg value: The vehicle operates more efficiently, consumes less fuel, and thus produces less CO2 emissions.
- Low mpg value: The vehicle consumes more fuel and produces more CO2 emissions.
 - Therefore, there is a negative relationship between “Fuel Consumption Comb (mpg)” and “CO2 Emissions.”
 - As fuel efficiency increases (mpg value increases), CO2 emissions decrease.

- This explains why environmentally friendly vehicles have high mpg values and produce fewer CO2 emissions.

Model

The “Model” column includes terms that identify specific features or configurations of vehicles:

- 4WD/4X4: Four-wheel drive. A drive system where all four wheels receive power.
- AWD: All-wheel drive. Similar to 4WD but often with more complex mechanisms for power distribution.
- FFV: Flexible-fuel vehicle. Vehicles that can use multiple types of fuel, such as both gasoline and ethanol blends.
- SWB: Short wheelbase.
- LWB: Long wheelbase.
- EWB: Extended wheelbase.

Transmission

The “Transmission” column indicates the type of transmission system in the vehicle:

- A: Automatic. A transmission type that operates without the need for the driver to manually change gears.
- AM: Automated manual. A version of a manual transmission that is automated.
- AS: Automatic with select shift. An automatic transmission that allows for manual intervention.
- AV: Continuously variable. A transmission that uses continuously varying ratios instead of fixed gear ratios.
- M: Manual. A transmission type that requires the driver to manually change gears.
- 3 – 10: Number of gears in the transmission.

Fuel Type

The “Fuel Type” column specifies the type of fuel used by the vehicle:

- X: Regular gasoline.
- Z: Premium gasoline.
- D: Diesel.
- E: Ethanol (E85).
- N: Natural gas.

Vehicle Class

The “Vehicle Class” column categorizes vehicles by size and type:

- COMPACT: Smaller-sized vehicles.
- SUV – SMALL: Smaller-sized sports utility vehicles.
- MID-SIZE: Medium-sized vehicles.
- TWO-SEATER: Vehicles with two seats.
- MINICOMPACT: Very small-sized vehicles.
- SUBCOMPACT: Smaller than compact-sized vehicles.
- FULL-SIZE: Larger-sized vehicles.
- STATION WAGON – SMALL: Smaller-sized station wagons.
- SUV – STANDARD: Standard-sized sports utility vehicles.
- VAN – CARGO: Vans designed for cargo.
- VAN – PASSENGER: Vans designed for passenger transportation.
- PICKUP TRUCK – STANDARD: Standard-sized pickup trucks.
- MINIVAN: Smaller-sized vans.
- SPECIAL PURPOSE VEHICLE: Vehicles designed for special purposes.
- STATION WAGON – MID-SIZE: Mid-sized station wagons.
- PICKUP TRUCK – SMALL: Smaller-sized pickup trucks.

This dataset can be used to understand the fuel efficiency and environmental impact of vehicles. Machine learning models can use these features to predict CO2 emissions or perform analyses comparing the fuel consumption of different vehicles.

EXPLORATORY DATA ANALYSIS(EDA)

```
[261]: import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import seaborn as sns

import plotly.express as px
import scipy.stats as stats
```

```
%matplotlib inline

from scipy import stats
from sklearn.preprocessing import PolynomialFeatures
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, MinMaxScaler, RobustScaler

from sklearn.linear_model import LinearRegression

from sklearn.model_selection import cross_val_score, cross_validate

from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error

from yellowbrick.regressor import ResidualsPlot, PredictionError

import warnings
warnings.filterwarnings("ignore")
```

```
[262]: co2 = pd.read_csv('CO2 Emissions_Canada.csv')
df = co2.copy()
```

Understanding the Data

```
[263]: df.head()
```

```
[263]:
```

	Make	Model	Vehicle Class	Engine Size(L)	Cylinders	Transmission	\
0	ACURA	ILX	COMPACT	2.0	4	AS5	
1	ACURA	ILX	COMPACT	2.4	4	M6	
2	ACURA	ILX HYBRID	COMPACT	1.5	4	AV7	
3	ACURA	MDX 4WD	SUV - SMALL	3.5	6	AS6	
4	ACURA	RDx AWD	SUV - SMALL	3.5	6	AS6	

	Fuel Type	Fuel Consumption City (L/100 km)	\
0	Z	9.9	
1	Z	11.2	
2	Z	6.0	
3	Z	12.7	
4	Z	12.1	

	Fuel Consumption Hwy (L/100 km)	Fuel Consumption Comb (L/100 km)	\
0	6.7	8.5	
1	7.7	9.6	
2	5.8	5.9	
3	9.1	11.1	
4	8.7	10.6	

	Fuel Consumption Comb (mpg)	CO2 Emissions(g/km)
0	33	196
1	29	221
2	48	136
3	25	255
4	27	244

[264]: *# Display random 5 sample*

```
df.sample(5)
```

[264]:

	Make	Model	Vehicle Class	Engine Size(L)	\
6688	CHEVROLET	Colorado 4WD	PICKUP TRUCK - SMALL	2.5	
6	ACURA	TL	MID-SIZE	3.5	
327	DODGE	GRAND CARAVAN FFV	MINIVAN	3.6	
369	FORD	F-150	PICKUP TRUCK - STANDARD	3.5	
6464	VOLKSWAGEN	Atlas	SUV - SMALL	2.0	

	Cylinders	Transmission	Fuel Type	Fuel Consumption City (L/100 km)	\
6688	4	A6	X	12.6	
6	6	AS6	Z	11.8	
327	6	A6	E	19.2	
369	6	A6	X	14.5	
6464	4	AS8	X	11.6	

	Fuel Consumption Hwy (L/100 km)	Fuel Consumption Comb (L/100 km)	\
6688	9.7	11.3	
6	8.1	10.1	
327	13.1	16.5	
369	10.6	12.7	
6464	9.1	10.5	

	Fuel Consumption Comb (mpg)	CO2 Emissions(g/km)
6688	25	265
6	28	232
327	17	264
369	22	292
6464	27	245

[265]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 7385 entries, 0 to 7384
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	Make	7385 non-null	object

```

1  Model                                7385 non-null  object
2  Vehicle Class                       7385 non-null  object
3  Engine Size(L)                     7385 non-null  float64
4  Cylinders                          7385 non-null  int64
5  Transmission                       7385 non-null  object
6  Fuel Type                          7385 non-null  object
7  Fuel Consumption City (L/100 km)   7385 non-null  float64
8  Fuel Consumption Hwy (L/100 km)    7385 non-null  float64
9  Fuel Consumption Comb (L/100 km)   7385 non-null  float64
10 Fuel Consumption Comb (mpg)         7385 non-null  int64
11 CO2 Emissions(g/km)                7385 non-null  int64
dtypes: float64(4), int64(3), object(5)
memory usage: 692.5+ KB

```

```
[266]: df.isnull().values.any()
```

```
[266]: False
```

```

[267]: # Check out the missing values

missing_count = df.isnull().sum()
value_count = df.isnull().count()
missing_percentage = round(missing_count / value_count * 100, 2)
missing_df = pd.DataFrame({"count": missing_count, "percentage":
    ↪missing_percentage})
missing_df

```

```

[267]:

```

	count	percentage
Make	0	0.0
Model	0	0.0
Vehicle Class	0	0.0
Engine Size(L)	0	0.0
Cylinders	0	0.0
Transmission	0	0.0
Fuel Type	0	0.0
Fuel Consumption City (L/100 km)	0	0.0
Fuel Consumption Hwy (L/100 km)	0	0.0
Fuel Consumption Comb (L/100 km)	0	0.0
Fuel Consumption Comb (mpg)	0	0.0
CO2 Emissions(g/km)	0	0.0

```

[268]: # Check out the duplicated values!!!!!!!

df.duplicated().sum()

```

```
[268]: 1103
```

```
[269]: duplicated_rows = df[df.duplicated(keep=False)]
```

```
duplicated_rows
```

```
[269]:
```

	Make	Model	Vehicle Class	Engine Size(L)	\
4	ACURA	RDX AWD	SUV - SMALL	3.5	
5	ACURA	RLX	MID-SIZE	3.5	
12	ALFA ROMEO	4C	TWO-SEATER	1.8	
13	ASTON MARTIN	DB9	MINICOMPACT	5.9	
15	ASTON MARTIN	V8 VANTAGE	TWO-SEATER	4.7	
...	
7356	TOYOTA	Tundra	PICKUP TRUCK - STANDARD	5.7	
7365	VOLKSWAGEN	Golf GTI	COMPACT	2.0	
7366	VOLKSWAGEN	Jetta	COMPACT	1.4	
7367	VOLKSWAGEN	Jetta	COMPACT	1.4	
7368	VOLKSWAGEN	Jetta GLI	COMPACT	2.0	

	Cylinders	Transmission	Fuel Type	Fuel Consumption City (L/100 km)	\
4	6	AS6	Z	12.1	
5	6	AS6	Z	11.9	
12	4	AM6	Z	9.7	
13	12	A6	Z	18.0	
15	8	AM7	Z	17.4	
...	
7356	8	AS6	X	17.7	
7365	4	M6	X	9.8	
7366	4	AS8	X	7.8	
7367	4	M6	X	7.9	
7368	4	AM7	X	9.3	

	Fuel Consumption Hwy (L/100 km)	Fuel Consumption Comb (L/100 km)	\
4	8.7	10.6	
5	7.7	10.0	
12	6.9	8.4	
13	12.6	15.6	
15	11.3	14.7	
...	
7356	13.6	15.9	
7365	7.3	8.7	
7366	5.9	7.0	
7367	5.9	7.0	
7368	7.2	8.4	

	Fuel Consumption Comb (mpg)	CO2 Emissions(g/km)
4	27	244
5	28	230
12	34	193

13	18	359
15	19	338
...
7356	18	371
7365	32	203
7366	40	162
7367	40	163
7368	34	196

[2102 rows x 12 columns]

Some duplicate rows are identical, while others have slight variations. Possible reasons for duplicates include:

1. **Data Entry Errors:** Same data entered multiple times.
2. **Different Periods:** Performance of the same model recorded over different years.
3. **Model Updates:** Comparing different versions of the same model.

Understanding the context and data collection methods is crucial to determine the cause of duplicates. Whether to drop duplicates depends on the analysis purpose:

- For unique observations, duplicates can be dropped.
- For analyzing changes over time or variations, keeping duplicates might be more appropriate.
- However, since I will be using a linear model, I do not plan to delete these rows because duplicate rows contain similar values. Linear regression tries to find the general trend of the data points and the fact that duplicate rows contain similar values does not lead to a major change in the model's predictions.

```
[270]: # Let's observe unique values

def get_unique_values(df):

    output_data = []

    for col in df.columns:

        # If the number of unique values in the column is less than or equal to 5
        if df.loc[:, col].nunique() <= 10:
            # Get the unique values in the column
            unique_values = df.loc[:, col].unique()
            # Append the column name, number of unique values, unique values,
            and data type to the output data
            output_data.append([col, df.loc[:, col].nunique(), unique_values,
            df.loc[:, col].dtype])
        else:
            # Otherwise, append only the column name, number of unique values,
            and data type to the output data
```



```

        output_data.append([col, df.loc[:, col].nunique(),"-", df.loc[:, col].dtype])

    output_df = pd.DataFrame(output_data, columns=['Column Name', 'Number of Unique Values', ' Unique Values ', 'Data Type'])
    #print(output_data)
    return output_df

```

```
[271]: get_unique_values(df)
```

```
[271]:
```

	Column Name	Number of Unique Values \
0	Make	42
1	Model	2053
2	Vehicle Class	16
3	Engine Size(L)	51
4	Cylinders	8
5	Transmission	27
6	Fuel Type	5
7	Fuel Consumption City (L/100 km)	211
8	Fuel Consumption Hwy (L/100 km)	143
9	Fuel Consumption Comb (L/100 km)	181
10	Fuel Consumption Comb (mpg)	54
11	CO2 Emissions(g/km)	331

	Unique Values	Data Type
0	-	object
1	-	object
2	-	object
3	-	float64
4	[4, 6, 12, 8, 10, 3, 5, 16]	int64
5	-	object
6	[Z, D, X, E, N]	object
7	-	float64
8	-	float64
9	-	float64
10	-	int64
11	-	int64

```
[272]: # Basic statistics summary of Numerical features
```

```
df.describe().T
```

```
[272]:
```

	count	mean	std	min	25% \
Engine Size(L)	7385.0	3.160068	1.354170	0.9	2.0
Cylinders	7385.0	5.615030	1.828307	3.0	4.0
Fuel Consumption City (L/100 km)	7385.0	12.556534	3.500274	4.2	10.1
Fuel Consumption Hwy (L/100 km)	7385.0	9.041706	2.224456	4.0	7.5

Fuel Consumption Comb (L/100 km)	7385.0	10.975071	2.892506	4.1	8.9
Fuel Consumption Comb (mpg)	7385.0	27.481652	7.231879	11.0	22.0
CO2 Emissions(g/km)	7385.0	250.584699	58.512679	96.0	208.0

	50%	75%	max
Engine Size(L)	3.0	3.7	8.4
Cylinders	6.0	6.0	16.0
Fuel Consumption City (L/100 km)	12.1	14.6	30.6
Fuel Consumption Hwy (L/100 km)	8.7	10.2	20.6
Fuel Consumption Comb (L/100 km)	10.6	12.6	26.1
Fuel Consumption Comb (mpg)	27.0	32.0	69.0
CO2 Emissions(g/km)	246.0	288.0	522.0

```
[273]: # Basic statistics summary of Object features
```

```
df.describe(include= 'object').T
```

```
[273]:
```

	count	unique	top	freq
Make	7385	42	FORD	628
Model	7385	2053	F-150 FFV	32
Vehicle Class	7385	16	SUV - SMALL	1217
Transmission	7385	27	AS6	1324
Fuel Type	7385	5	X	3637

```
[274]: df.columns
```

```
[274]: Index(['Make', 'Model', 'Vehicle Class', 'Engine Size(L)', 'Cylinders',
        'Transmission', 'Fuel Type', 'Fuel Consumption City (L/100 km)',
        'Fuel Consumption Hwy (L/100 km)', 'Fuel Consumption Comb (L/100 km)',
        'Fuel Consumption Comb (mpg)', 'CO2 Emissions(g/km)'],
        dtype='object')
```

Rename the Columns

```
[275]: df.rename(columns={ 'Make': 'make',
        'Model': 'model',
        'Vehicle Class': 'vehicle_class',
        'Engine Size(L)': 'engine_size',
        'Cylinders': 'cylinders',
        'Transmission': 'transmission',
        'Fuel Type': 'fuel_type',
        'Fuel Consumption City (L/100 km)': 'fuel_cons_city',
        'Fuel Consumption Hwy (L/100 km)': 'fuel_cons_hwy',
        'Fuel Consumption Comb (L/100 km)': 'fuel_cons_comb',
        'Fuel Consumption Comb (mpg)': 'fuel_cons_comb_mpg',
        'CO2 Emissions(g/km)': 'co2'
        }, inplace=True)
```

```
[276]: df.columns
```

```
[276]: Index(['make', 'model', 'vehicle_class', 'engine_size', 'cylinders',  
        'transmission', 'fuel_type', 'fuel_cons_city', 'fuel_cons_hwy',  
        'fuel_cons_comb', 'fuel_cons_comb_mpg', 'co2'],  
        dtype='object')
```

```
## Data Visualisation
```

```
## Categorical Features
```

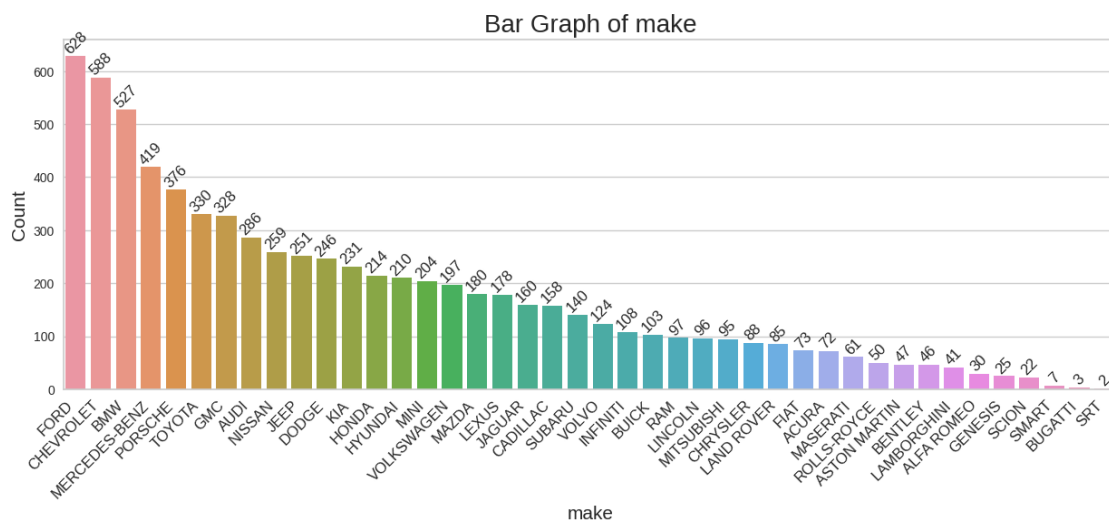
```
### Distribution of Categorical Features
```



```
[277]: # Let's look at the distribution of our categorical characteristics with a bar  
        ↪graph
```




```
def plot_bar_graphs(df, columns):  
    for column in columns:  
        plt.figure(figsize=(15, 5))  
        ax = sns.countplot(x=column, data=df, order=df[column].value_counts().  
        ↪index)  
        ax.bar_label(ax.containers[0], rotation=45)  
        plt.xlabel(column, fontsize=15)  
        plt.ylabel('Count', fontsize=15)  
        plt.title(f'Bar Graph of {column}', fontsize=20)  
        plt.xticks(rotation=45, ha='right', fontsize=12)  
        plt.show()
```

```
cat_features = ['make', 'vehicle_class', 'engine_size', 'cylinders',  
        ↪'transmission', 'fuel_type']
```

```
plot_bar_graphs(df, cat_features)
```








```
[278]: # Model Feature
```

```
df.model.unique()
```

```
[278]: array(['ILX', 'ILX HYBRID', 'MDX 4WD', ...,  
        'Tacoma 4WD D-Cab TRD Off-Road/Pro', 'Atlas Cross Sport 4MOTION',  
        'XC40 T4 AWD'], dtype=object)
```

```
[279]: fig = plt.figure(figsize = (15,10))  
ax = fig.add_axes([0,0,1,1])  
counts = df.model.value_counts().sort_values(ascending=False).head(20)  
counts.plot(kind = "bar")  
plt.title('Top 20 Model')  
plt.xlabel('model')  
plt.ylabel('Number of Cars')  
plt.xticks(rotation = 45)  
ax.bar_label(ax.containers[0], labels=counts.values, fontsize=12);
```



0.2 Conclusion:

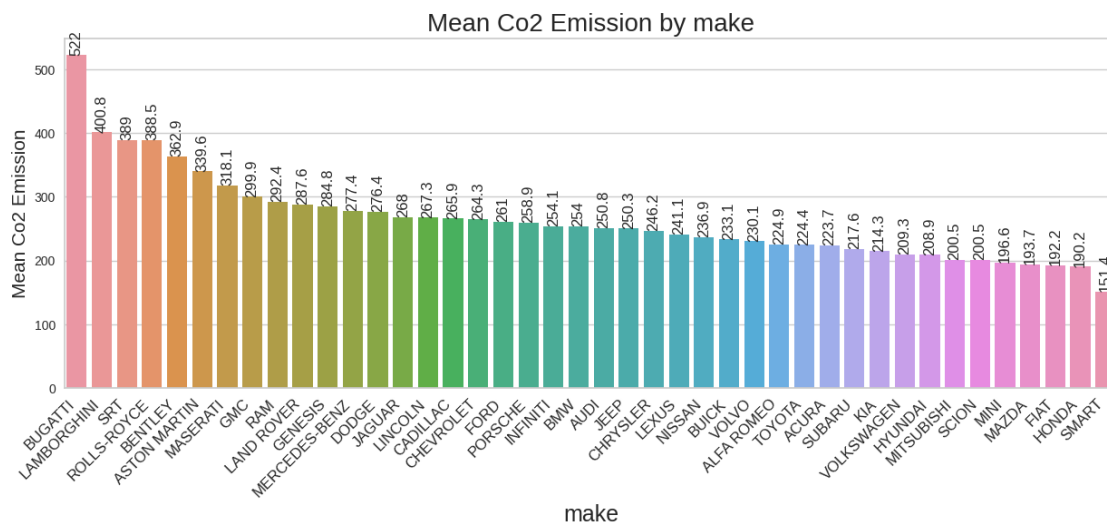
As can be seen from the graphs above: - The number of vehicles consuming diesel, ethanol and natural gas fuel in the data set is very small. - Widespread use of AS6, AS8, M6, A6, A9 as transmission options - 4, 6, 8 are commonly used as cylinders option - Engine Size (L) with 2.0 and 3.0 options in density - The dataset is generally dominated by smaller sized vehicles




Target Variable vs Categorical Features

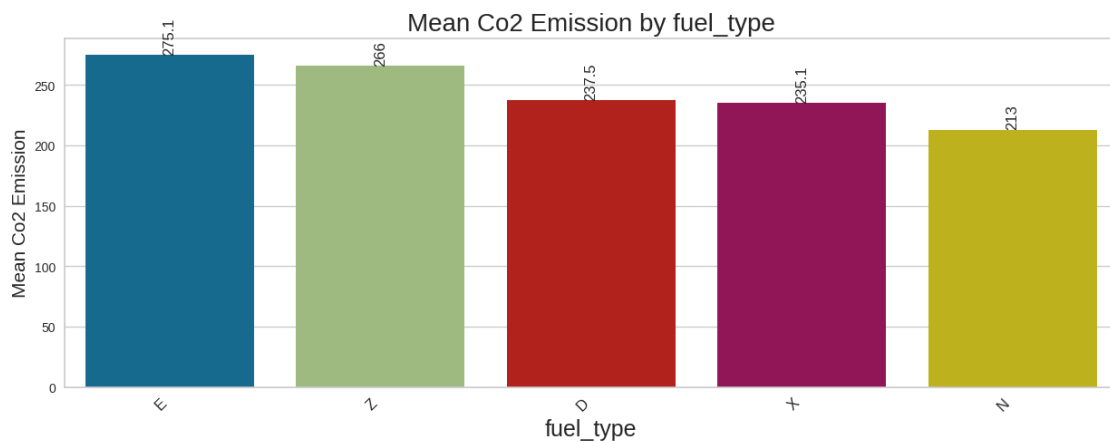

```
[280]: # Let's look at the relationship between our categorical attributes and the
        ↪target variable

def plot_bar_with_co2(df, columns):
    for column in columns:
        plt.figure(figsize=(15, 5))
        grouped_data = df.groupby(column)['co2'].mean().round(1).reset_index()
        grouped_data_sorted = grouped_data.sort_values(by='co2',
        ↪ascending=False)
        ax = sns.barplot(x=column, y='co2', data=grouped_data_sorted,
        ↪order=grouped_data_sorted[column])
        ax.bar_label(ax.containers[0], rotation=90)
        plt.xlabel(column, fontsize=18)
        plt.ylabel('Mean Co2 Emission', fontsize=15)
        plt.title(f'Mean Co2 Emission by {column}', fontsize=20)
        plt.xticks(rotation=45, ha='right', fontsize=12)
        plt.show()

plot_bar_with_co2(df, cat_features)
```







0.3 Conclusion:

As can be seen from the graphs above:

- Bugatti has the highest average Co2 emissions
- Large-volume vehicles have high CO2 emission averages
- CO2 emission averages of high volume and cylinders engines are also high
- Ethanol and Premium gasoline is the fuel with the highest average CO2

ANOVA Test for Categorical Features

```
[281]: # Perform ANOVA test for each categorical feature
anova_results = {}
categorical_features = df.select_dtypes(include=['object']).columns
```

```

for feature in categorical_features:
    groups = [df["co2"][df[feature] == category].values for category in
    ↪df[feature].unique()]
    anova_results[feature] = stats.f_oneway(*groups)

# Display the ANOVA results
for feature, result in anova_results.items():
    print(f"ANOVA result for {feature}:")
    print(f"F-statistic: {result.statistic}, p-value: {result.pvalue}")
    print()

```

ANOVA result for make:

F-statistic: 106.8000265413262, p-value: 0.0

ANOVA result for model:

F-statistic: 58.405091462865016, p-value: 0.0

ANOVA result for vehicle_class:

F-statistic: 266.0228094521597, p-value: 0.0

ANOVA result for transmission:

F-statistic: 103.70394951088048, p-value: 0.0

ANOVA result for fuel_type:

F-statistic: 148.94555963595639, p-value: 1.062810397301377e-122

NOTE:

- The p-values for each of the Make, Model, Vehicle Class, Transmission, and Fuel Type variables are much smaller than 0.05, indicating that these variables create statistically significant differences in co2_emissions.

Label Encoding the Categorical Features

Converting all categorical column into numerical category for correlation matrix

```

[282]: from sklearn.preprocessing import LabelEncoder

# Copy the original dataframe to avoid modifying it directly
df_labeled = df.copy()

# List of categorical columns
categorical_columns = ['make', 'model', 'vehicle_class', 'transmission',
    ↪'fuel_type']

# Apply Label Encoding to each categorical column
label_encoders = {}
for column in categorical_columns:

```

```

le = LabelEncoder()
df_labeled[column] = le.fit_transform(df_labeled[column])
label_encoders[column] = le

# Display the first few rows of the labeled dataframe
print(df_labeled.head())

```

	make	model	vehicle_class	engine_size	cylinders	transmission	\
0	0	1057	0	2.0	4	14	
1	0	1057	0	2.4	4	25	
2	0	1058	0	1.5	4	22	
3	0	1233	11	3.5	6	15	
4	0	1499	11	3.5	6	15	

	fuel_type	fuel_cons_city	fuel_cons_hwy	fuel_cons_comb	\
0	4	9.9	6.7	8.5	
1	4	11.2	7.7	9.6	
2	4	6.0	5.8	5.9	
3	4	12.7	9.1	11.1	
4	4	12.1	8.7	10.6	

	fuel_cons_comb_mpg	co2
0	33	196
1	29	221
2	48	136
3	25	255
4	27	244

[283]: df_labeled.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7385 entries, 0 to 7384
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   make                  7385 non-null  int64
1   model                 7385 non-null  int64
2   vehicle_class         7385 non-null  int64
3   engine_size           7385 non-null  float64
4   cylinders              7385 non-null  int64
5   transmission          7385 non-null  int64
6   fuel_type             7385 non-null  int64
7   fuel_cons_city         7385 non-null  float64
8   fuel_cons_hwy          7385 non-null  float64
9   fuel_cons_comb         7385 non-null  float64
10  fuel_cons_comb_mpg     7385 non-null  int64
11  co2                    7385 non-null  int64
dtypes: float64(4), int64(8)


```

memory usage: 692.5 KB

Correlations of Numerical Features

```
[284]: correlation_matrix = df_labeled.corr()

plt.figure(figsize=(10,8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', linewidths=0.5)
plt.title('Correlation Matrix Heatmap')
plt.show()
```



- The correlation matrix shows that features like **engine_size**, **cylinders**, and fuel consumption metrics (city, highway, combined) have strong positive correlations with **co2_emissions**.
- Notably, **fuel_cons_mpg** has a strong negative correlation with CO2 emissions, indicating that higher fuel efficiency results in lower emissions.
- From a multicollinearity perspective, **fuel_cons_city**, **fuel_cons_hwy**, and **fuel_cons_comb** are highly intercorrelated, suggesting potential redundancy.

- It may be beneficial to select just one or combine them into a single metric to avoid multicollinearity.
- Key predictors for a CO2 emissions model include **engine_size**, **cylinders**, and fuel consumption metrics.
- Categorical features such as **make**, **model**, **vehicle_class**, **transmission**, and **fuel_type** are also important due to their significant impact on CO2 emissions.
 - Managing multicollinearity among these highly correlated features is crucial to ensure model stability and performance.

Numerical Features


Distribution of Numerical Features


```
[285]: numerical_df = df.select_dtypes(include=['number'])


plt.figure(figsize=(15, 10))


num_vars = len(numerical_df.columns)


for i, var in enumerate(numerical_df.columns, 1):
    sns.histplot(data=df, x=var, kde=True, label = f'Distribution of_
↪{var}',color = "blue")
    plt.ylabel('Frequency', fontsize=10)
    plt.tight_layout()
    plt.show()
```


















```
[286]: # Consumption of Highway and City

plt.figure(figsize=(8, 5))
sns.histplot(data=df, x="fuel_cons_hwy", kde=True, label = "Fuel Consumption in Highway",color = "orange")
sns.histplot(data=df, x="fuel_cons_city", kde=True, label = "Fuel Consumption in City")
plt.xlabel('Consumption (L/100 km)', fontsize=10)
plt.ylabel('Frequency', fontsize=10)
plt.title(f'Histogram of Highway and City', fontsize=10)
plt.legend()
plt.show()
```



0.4 Observations

1. **Engine Size (engine_size):** The distribution is skewed with several peaks, indicating different types of vehicles with varying engine sizes.
2. **Cylinders:** The distribution is concentrated around specific values (4, 6, 8), reflecting common engine types in the dataset.
3. **Fuel Consumption City (fuel_cons_city):** The distribution is approximately normal, which is beneficial for model learning.
4. **Fuel Consumption Hwy (fuel_cons_hwy):** Similar to city fuel consumption, this also shows an approximately normal distribution.
5. **Fuel Consumption Combined (fuel_cons_comb):** The combined fuel consumption distribution is normal, making it a useful variable for modeling.
6. **Fuel Consumption MPG (fuel_cons_mpg):** This shows a normal distribution and has an inverse relationship with other fuel consumption metrics, which is expected.
7. **CO2 Emissions (co2):** The distribution is nearly normal, which is advantageous for regression models.


Overall Evaluation: - Most features exhibit normal or skew-normal distributions, which are suitable for modeling. - Consider multicollinearity, especially among `fuel_cons_City`, `fuel_cons_Hwy`, and `fuel_cons_Comb`. - Features like `cylinders`, `engine_size`, and `fuel_cons_mpg` provide important information about vehicle performance and efficiency, making them valuable for the model.

Target Variable vs Numerical Features

Correlations of Numerical Features

```
[287]: correlation_matrix = df.corr(numeric_only=True)

plt.figure(figsize=(8,6))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', linewidths=0.5)
plt.title('Correlation Matrix Heatmap')
plt.show()
```



The correlation matrix shows strong multicollinearity among several features. Specifically:

- **fuel_cons_city**, **fuel_cons_hwy**, and **fuel_cons_comb** have very high correlations (0.95 and above), indicating they carry almost identical information.
- **engine_size** and **cylinders** are also highly correlated (0.93).
- **co2_emissions** has strong positive correlations with **fuel_cons_hwy**, **fuel_cons_comb**, **fuel_cons_city**, and **engine_size** (0.85 and above), indicating these features are important for predicting CO2 emissions.
- **fuel_cons_mpg** shows high negative correlations with other fuel consumption measures (-0.93 and above), as higher mpg indicates lower fuel consumption.

These findings suggest that removing some highly correlated features can help reduce multicollinear-

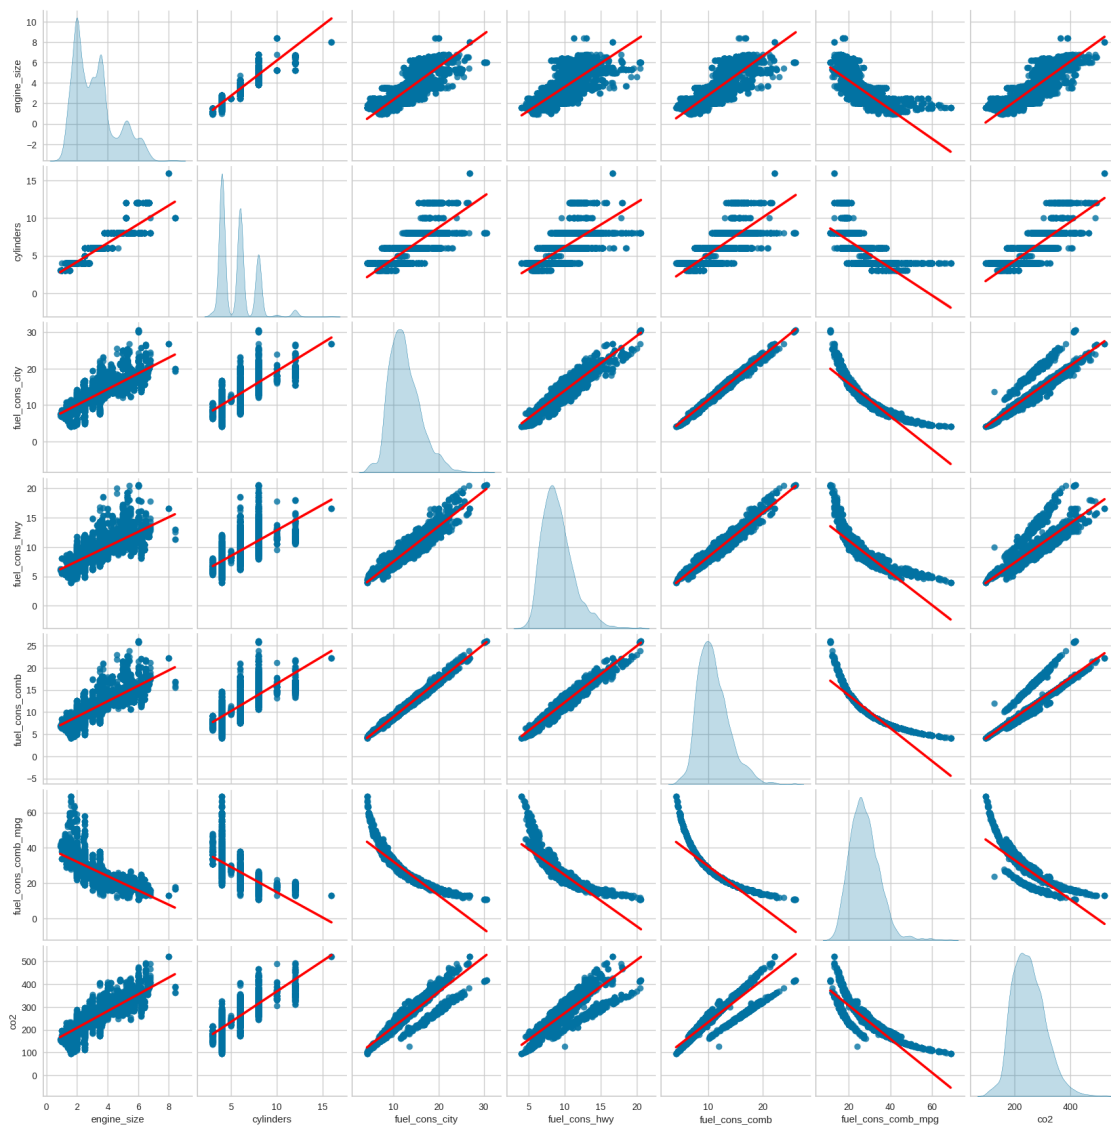
ity and improve model performance.

0.5 Pairplot for the Numerical Data of dataframe

```
[288]: # Pairplot for the dataframe
```

```
sns.pairplot(df,  
              kind="reg",  
              diag_kind="kde",  
              plot_kws={"line_kws": {"color": "red"}}  
              )
```

```
[288]: <seaborn.axisgrid.PairGrid at 0x7f23246dbd10>
```



The pairplot shows the relationships and distributions between various numerical features in the dataset. Here's a brief analysis:

1. **Engine Size (engine_size):**
 - Positively correlated with cylinders, fuel_cons_city, fuel_cons_hwy, fuel_cons_comb, and co2.
 - Larger engines tend to have more cylinders and higher fuel consumption.
2. **Cylinders:**
 - Strong positive correlation with engine_size and fuel consumption metrics.
 - As the number of cylinders increases, fuel consumption and CO2 emissions also increase.
3. **Fuel Consumption City (fuel_cons_City):**
 - High positive correlation with fuel_cons_hwy, fuel_cons_comb, and co2.
 - Vehicles that consume more fuel in the city tend to consume more on highways and produce higher CO2 emissions.
4. **Fuel Consumption Hwy (fuel_cons_Hwy):**
 - Similar correlations as city fuel consumption, showing strong positive relationships with fuel_cons_city, fuel_cons_comb, and co2.
5. **Fuel Consumption Combined (fuel_cons_Comb):**
 - Very high correlation with both city and highway fuel consumption.
 - Indicative of overall vehicle efficiency.
6. **Fuel Consumption MPG (fuel_cons_mpg):**
 - Shows a strong negative correlation with other fuel consumption metrics and co2.
 - *Higher MPG values indicate better fuel efficiency and lower CO2 emissions.*
 - fuel efficiency increases (mpg value increases), CO2 emissions decrease.
7. **CO2 Emissions (co2):**
 - Strongly correlated with engine_size, cylinders, fuel_cons_city, fuel_cons_hwy, and fuel_cons_comb.
 - *Vehicles with larger engines, more cylinders, and higher fuel consumption emit more CO2.*

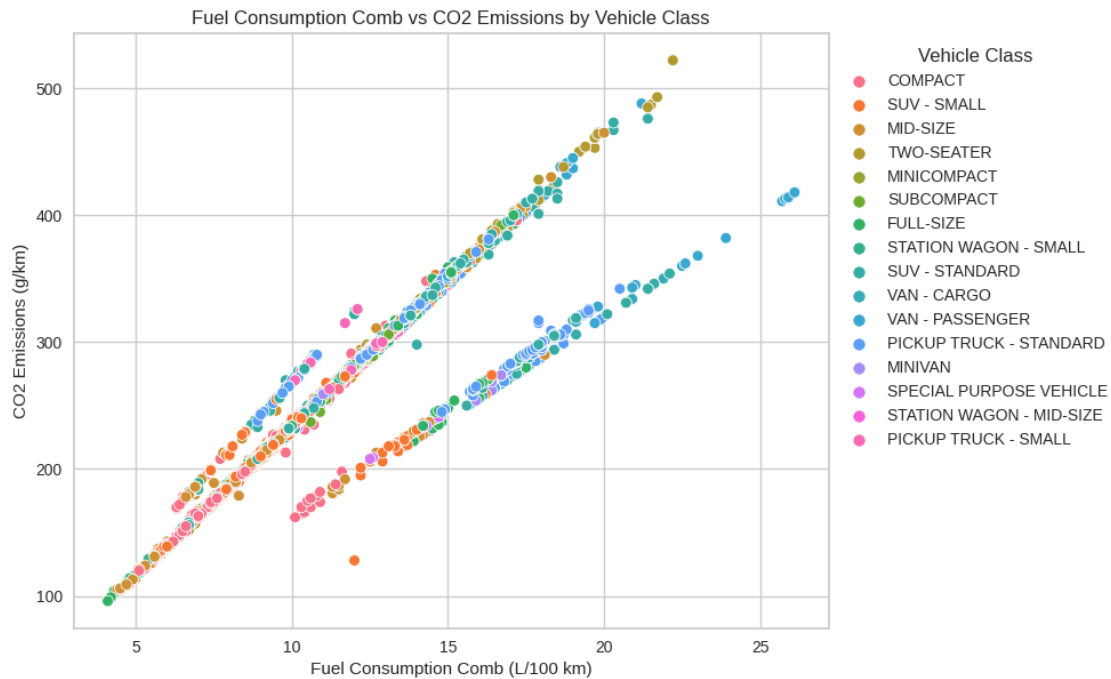
Overall Evaluation: - The pairplot reveals strong relationships among features, particularly between engine size, fuel consumption, and CO2 emissions. - The negative correlation between MPG and other features highlights its importance in representing fuel efficiency. - Suggest focusing on features like fuel_cons_comb, engine_size, and fuel_cons_mpg for predictive modeling, while considering multicollinearity.

```
[289]: # Target vs Fuel Consumption Combined (city+hwy)
# Hue: Vehicle Class

plt.figure(figsize=(10,6))
sns.scatterplot(data=df,x='fuel_cons_comb',y='co2',hue='vehicle_class')
plt.legend(bbox_to_anchor=(1, 1), loc='upper left', title='Vehicle Class')
plt.tight_layout()

plt.title('Fuel Consumption Comb vs CO2 Emissions by Vehicle Class')
plt.xlabel('Fuel Consumption Comb (L/100 km)')
plt.ylabel('CO2 Emissions (g/km)')
```

[289]: Text(86.47222222222221, 0.5, 'CO2 Emissions (g/km)')



0.5.1 Observation:


- Vehicle class makes impact on overall fuel efficiency of vehicle as well lead to more emission
- Two-Seater, Mid Size, Passenger van, Cargo van, Pick Up Truck have most emission with lowest fuel efficiency
- Mini Van, SPV, Station Wagon are more fuel efficient with low emission

```
[290]: # Target vs Fuel Consumption in Miles Per Gallon (mpg)
# Hue: Fuel Type

plt.figure(figsize=(8,5))
sns.scatterplot(data=df,x='fuel_cons_comb_mpg',y='co2',hue='fuel_type')

plt.title('Fuel Consumption mpg(efficiency) vs CO2 Emissions by Fuel_Type')
plt.xlabel('Fuel Consumption in Miles Per Gallon (L/100 km)')
plt.ylabel('CO2 Emissions (g/km)')
```

[290]: Text(0, 0.5, 'CO2 Emissions (g/km)')



0.5.2 Observation:

- Fuel Type makes impact on overall fuel efficiency of vehicle as well lead to less emission with Fuel Consumption
- Premium gasolin and Ethenol have most emission with lowest fuel efficiency
- Regular gasoline has more efficiency of fuel as well lead to less emission
- Dieselis is more fuel efficienct with low emission

Outlier Analysis


```
[291]: # Checking Outliers

# Initialize the subplot counter
x = 0

# Create a figure with specified size
plt.figure(figsize=(16, 4))

# Loop through each numerical column and create a boxplot
for col in df.select_dtypes(include=['number']).columns:
    x += 1
    plt.subplot(1, 8, x)
    sns.boxplot(data=df[col])
    plt.title(col)
```

```
# Show the plots
plt.tight_layout() # Adjust subplots to fit in the figure area.
plt.show()
```



Skewness

- Calculate skewness for numeric features
- A skewness value greater than 1 indicates positive skewness,
- a skewness value less than -1 indicates negative skewness,
- and a skewness value close to zero indicates a relatively symmetric distribution.

```
[292]: num_cols= df.select_dtypes('number').columns

skew_limit = 0.75 # define a limit above which we will log
↳ transform
skew_vals = df[num_cols].skew()

# Showing the skewed columns
skew_cols = (skew_vals
              .sort_values(ascending=False)
              .to_frame()
              .rename(columns={0: 'Skew'})
              .query('abs(Skew) > {}'.format(skew_limit)))

skew_cols
```

```
[292]:
```

	Skew
cylinders	1.110415
fuel_cons_hwy	1.079217
fuel_cons_comb_mpg	0.977034
fuel_cons_comb	0.893316
engine_size	0.809181
fuel_cons_city	0.809005

MACHINE LEARNING

- The objective of creating and using a model with the CO2 emission dataset is to build machine learning algorithms capable of accurately predicting vehicle CO2 emissions based on their characteristics.
- By examining variables such as engine size, number of cylinders, and fuel consumption, the aim is to develop models that can evaluate the environmental impact of different vehicles and guide policy decisions aimed at reducing carbon emissions.
- Additionally, these models can support automotive manufacturers in designing more fuel-efficient vehicles and help consumers make informed choices when selecting vehicles with lower carbon footprints.
- Ultimately, the goal is to harness data-driven insights to mitigate the environmental impact of transportation and promote sustainable development.
- Evaluating model accuracy on both training and test sets is essential to determine whether the model is overfitting or underfitting the data, addressing the bias-variance tradeoff effectively.


Simple Linear Regression Model

- This simple linear regression model was built using only `fuel_cons_comb` as the predictor and the target variable `co2`, without any data manipulation. This is often referred to as a “vanilla model.”

```
[293]: # Check the correlation between independent feature (fuel_cons_comb) and target_
        ↪variable (co2_emissions)

sns.regplot(x = 'fuel_cons_comb', y = 'co2', data=df, ci=None)
```

```
[293]: <Axes: xlabel='fuel_cons_comb', ylabel='co2'>
```



Splitting the Data

```
[294]: # Split the selected independent feature (fuel_cons_comb) and target variable
        ↪ (co2_emissions) for SIMPLE Linear Regression

X = df[['fuel_cons_comb']]
y = df['co2']
```

Train | Test Split

```
[295]: from sklearn.model_selection import train_test_split, cross_val_score,
        ↪ cross_val_predict
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
        ↪ random_state=42)
```

```
[296]: # Display the shapes of the resulting datasets
```

```
print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)
```

```
X_train shape: (5538, 1)
X_test shape: (1847, 1)
y_train shape: (5538,)
y_test shape: (1847,)
```

```
### Model
```

```
[297]: from sklearn.linear_model import LinearRegression
lin_reg = LinearRegression()
```

```
### Training the Model
```

```
[298]: lin_reg.fit(X_train, y_train)
```

```
[298]: LinearRegression()
```

```
### Predicting Test Data
```

```
[299]: # Predict using the model on the test data

y_pred_test = lin_reg.predict(X_test)
y_pred_train = lin_reg.predict(X_train)
```

```
### Evaluating the Model
```

```
[300]: # Comparing Actual y_test, Predicted y_test and Residuals
my_dict = {"Actual": y_test, "pred": y_pred_test, "residual": y_test -
↪ y_pred_test}
compare = pd.DataFrame(my_dict)
compare.head(20)
```

```
[300]:
```

	Actual	pred	residual
7261	253	249.309805	3.690195
4489	344	320.052022	23.947978
1539	322	307.020561	14.979439
3532	297	282.819276	14.180724
6418	308	292.127463	15.872537
3703	406	368.454592	37.545408
5976	242	240.001619	1.998381
4332	216	219.523608	-3.523608
5015	246	241.863256	4.136744
2087	223	226.970158	-3.970158
2126	283	275.372727	7.627273
4161	326	303.297287	22.702713
4814	274	264.202903	9.797097
486	251	249.309805	1.690195
6607	322	303.297287	18.702713
1128	382	355.423131	26.576869
5159	248	243.724893	4.275107

5391	193	200.907235	-7.907235
6643	204	208.353785	-4.353785
6003	211	213.938696	-2.938696

Performance Metrics

```
[301]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

R2Score_test=r2_score(y_test, y_pred_test)
mae_test=mean_absolute_error(y_test, y_pred_test)
mse_test=mean_squared_error(y_test, y_pred_test)
rmse_test=np.sqrt(mean_squared_error(y_test, y_pred_test))

R2Score_train=r2_score(y_train, y_pred_train)
mae_train=mean_absolute_error(y_train, y_pred_train)
mse_train=mean_squared_error(y_train, y_pred_train)
rmse_train=np.sqrt(mean_squared_error(y_train, y_pred_train))

print("performance Metrics of Simple linear Regression for Train and Test data")
print("\t\t Train_data  | Test_data")
print("-----")
print("R2 Score ",R2Score_train,"|",R2Score_test)
print("MAE      ",mae_train,"|",mae_test)
print("MSE      ",mse_train,"|",mse_test)
print("RMSE     ",rmse_train,"|",rmse_test)

scores = {
    "Simple_test": {"R2" : R2Score_test,
                    "mae" : mae_test,
                    "mse" : mse_test,
                    "rmse" : rmse_test},

    "Simple_train": {"R2" : R2Score_train,
                     "mae" : mae_train,
                     "mse" : mse_train,
                     "rmse" : rmse_train}
}

slr_score=pd.DataFrame(scores)

slr_score
```

```
performance Metrics of Simple linear Regression for Train and Test data
                Train_data  | Test_data
-----
R2 Score  0.8443308529162576 | 0.8381501636019864
```

MAE	13.901692049585145		14.233495693123851
MSE	533.4385730402728		552.1010877577722
RMSE	23.096289161687267		23.496831440808613

```
[301]:      Simple_test  Simple_train
R2      0.838150      0.844331
mae     14.233496      13.901692
mse     552.101088      533.438573
rmse    23.496831      23.096289
```

```
[302]: slr_score
```

```
[302]:      Simple_test  Simple_train
R2      0.838150      0.844331
mae     14.233496      13.901692
mse     552.101088      533.438573
rmse    23.496831      23.096289
```

```
[303]: rmse_test/df['co2'].mean()
```

```
[303]: 0.0937680215968133
```

- To determine how much the error deviates from the mean of the target label.
- According to the RMSE metric, our model has an average error rate of 9.3%.
- Prefer the RMSE metric because it penalizes poor predictions.

Conclusion - This simple linear regression model was built using only Fuel Consumption Combined (city+hwy) as the predictor and the target variable, without any data manipulation.

- The model explains 83.8% of the variance in the target variable, indicating a good fit but with room for improvement.

Multiple Linear Regression Model

- We will now create a multiple linear regression model using 'engine_size', 'fuel_cons_comb', 'fuel_cons_hwy', and 'fuel_cons_city' as the independent variables and the target variable.
- This model aims to capture the relationship between multiple predictors and the target variable for better prediction accuracy and insights.

Splitting the Data

```
[304]: X = df[["engine_size", "cylinders", "fuel_cons_comb", "fuel_cons_comb_mpg"]]
y = df["co2"]
```

```
[305]: X.head()
```

```
[305]:   engine_size  cylinders  fuel_cons_comb  fuel_cons_comb_mpg
0          2.0          4             8.5             33
1          2.4          4             9.6             29
2          1.5          4             5.9             48
```

3	3.5	6	11.1	25
4	3.5	6	10.6	27

```
[306]: # Check Multicollinarty between features
pd.DataFrame(X).corr()
```

```
[306]:
```

	engine_size	cylinders	fuel_cons_comb	fuel_cons_comb_mpg
engine_size	1.000000	0.927653	0.817060	-0.757854
cylinders	0.927653	1.000000	0.780534	-0.719321
fuel_cons_comb	0.817060	0.780534	1.000000	-0.925576
fuel_cons_comb_mpg	-0.757854	-0.719321	-0.925576	1.000000

Train | Test Split

```
[307]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↳random_state=42)
```

```
[308]: print("X Train features shape: {}\ny Train features shape: {}\nX Test features_
↳shape: {}\ny Test features shape: {}".
        .format(X_train.shape, y_train.shape, X_test.shape, y_test.shape))
```

```
X Train features shape: (5169, 4)
y Train features shape: (5169,)
X Test features shape: (2216, 4)
y Test features shape: (2216,)
```

Model

```
[309]: from sklearn.linear_model import LinearRegression

Multi_lin_reg = LinearRegression()
```

Training the Model

```
[310]: Multi_lin_reg.fit(X_train, y_train)
```

```
[310]: LinearRegression()
```

Predicting Test Data

```
[311]: y_pred_test = Multi_lin_reg.predict(X_test)
y_pred_train = Multi_lin_reg.predict(X_train)
```

```
[312]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

R2Score_test=r2_score(y_test, y_pred_test)
mae_test=mean_absolute_error(y_test, y_pred_test)
```



```

mse_test=mean_squared_error(y_test, y_pred_test)
rmse_test=np.sqrt(mean_squared_error(y_test, y_pred_test))

R2Score_train=r2_score(y_train, y_pred_train)
mae_train=mean_absolute_error(y_train, y_pred_train)
mse_train=mean_squared_error(y_train, y_pred_train)
rmse_train=np.sqrt(mean_squared_error(y_train, y_pred_train))

print("performance Metrics of Multiple Linear Regression for Train and Test_
↳data")
print("\t\t Train_data   | Test_data")
print("-----")
print("R2 Score ",R2Score_train,"|",R2Score_test)
print("MAE      ",mae_train,"|",mae_test)
print("MSE      ",mse_train,"|",mse_test)
print("RMSE     ",rmse_train,"|",rmse_test)

scores = {
    "MultiLinear_test": {"R2" : R2Score_test,
        "mae" : mae_test,
        "mse" : mse_test,
        "rmse" : rmse_test},

    "MultiLinear_train": {"R2" : R2Score_train,
        "mae" : mae_train,
        "mse" : mse_train,
        "rmse" : rmse_train}
}

Mlr_score=pd.DataFrame(scores)

Mlr_score

```

```

performance Metrics of Multiple Linear Regression for Train and Test data
                Train_data   | Test_data
-----
R2 Score  0.9038767582625543 | 0.9001077496796939
MAE       11.43668932277986  | 11.514613831965189
MSE       330.8291640651209  | 337.4066070610968
RMSE      18.18870979660517  | 18.368631061162308

```

```

[312]:      MultiLinear_test  MultiLinear_train
R2          0.900108      0.903877
mae         11.514614     11.436689
mse         337.406607     330.829164

```

rmse 18.368631 18.188710

Cross Validation for Multiple Linear

```
[313]: from sklearn.model_selection import cross_validate, cross_val_score

model = LinearRegression()
scores = cross_validate(model, X_train, y_train,
                        scoring = ['r2', \
↪ 'neg_mean_absolute_error', 'neg_mean_squared_error', \
                                'neg_root_mean_squared_error'], cv = 10, \
↪ return_train_score=True)
```

```
[314]: pd.DataFrame(scores, index = range(1,11))
```

```
[314]:
```

	fit_time	score_time	test_r2	train_r2	test_neg_mean_absolute_error \
1	0.004916	0.001419	0.917568	0.902321	-10.966634
2	0.002119	0.001454	0.898006	0.904529	-11.705880
3	0.001940	0.001363	0.892331	0.904956	-11.329495
4	0.002944	0.001299	0.920210	0.902001	-10.763197
5	0.001949	0.001306	0.875318	0.906981	-12.949277
6	0.002143	0.001394	0.894262	0.904821	-11.368600
7	0.002121	0.001284	0.915373	0.902485	-10.922760
8	0.001896	0.001249	0.917338	0.902337	-11.034622
9	0.002109	0.001174	0.898096	0.904445	-11.276306
10	0.004394	0.001843	0.902799	0.903967	-12.168696

	train_neg_mean_absolute_error	test_neg_mean_squared_error \
1	-11.495846	-280.218171
2	-11.445247	-358.809708
3	-11.403127	-322.546210
4	-11.491006	-279.055713
5	-11.241349	-428.067312
6	-11.380240	-339.357892
7	-11.508300	-310.967821
8	-11.543714	-287.793779
9	-11.467861	-326.238550
10	-11.382183	-383.955085

	train_neg_mean_squared_error	test_neg_root_mean_squared_error \
1	-336.492891	-16.739718
2	-327.759612	-18.942273
3	-331.783466	-17.959572
4	-336.650751	-16.704961
5	-320.045190	-20.689788
6	-329.958182	-18.421669
7	-333.053611	-17.634280

8	-335.691214	-16.964486
9	-331.372363	-18.062075
10	-325.032680	-19.594772

	train_neg_root_mean_squared_error
1	-18.343743
2	-18.104132
3	-18.214924
4	-18.348045
5	-17.889807
6	-18.164751
7	-18.249756
8	-18.321878
9	-18.203636
10	-18.028663

```
[315]: pd.DataFrame(scores, index = range(1,11)).iloc[:, 2:].mean()
```

```
[315]: test_r2          0.903130
train_r2          0.903884
test_neg_mean_absolute_error    -11.448547
train_neg_mean_absolute_error    -11.435887
test_neg_mean_squared_error     -331.701024
train_neg_mean_squared_error     -330.783996
test_neg_root_mean_squared_error  -18.171359
train_neg_root_mean_squared_error -18.186934
dtype: float64
```

- The fact that this score obtained after Cross Validation and Train-test score are compatible indicates that the model has generalization ability.

Compatibility and Generalization Ability:

- **Good Generalization:** If both the cross-validation score and the train-test score are similar and high, it indicates that the model is performing well both on the training data and the unseen test data. This consistency suggests that the model is not overfitting to the training data and is likely to generalize well to new, unseen data.
- **Overfitting:** Conversely, if the cross-validation score is high but the train-test score is significantly lower, this might indicate that the model is overfitting to the training data and may not perform as well on new, unseen data.
- **Underfitting:** If both scores are low, it may suggest that the model is underfitting and not capturing the underlying patterns in the data effectively.

Comparing the Scores Multi & Simple Linear Regression

```
[316]: pd.concat([slr_score, Mlr_score], axis=1)

# Concatinated simple linear ve multiple linear models scores
```

```
[316]:
```

	Simple_test	Simple_train	MultiLinear_test	MultiLinear_train
R2	0.838150	0.844331	0.900108	0.903877
mae	14.233496	13.901692	11.514614	11.436689
mse	552.101088	533.438573	337.406607	330.829164
rmse	23.496831	23.096289	18.368631	18.188710

```
[317]: rmse_test/df['co2'].mean()
```

```
[317]: 0.07330308337044823
```

```
[333]: # Create a figure and axis for the plot
plt.figure(figsize=(10, 6))

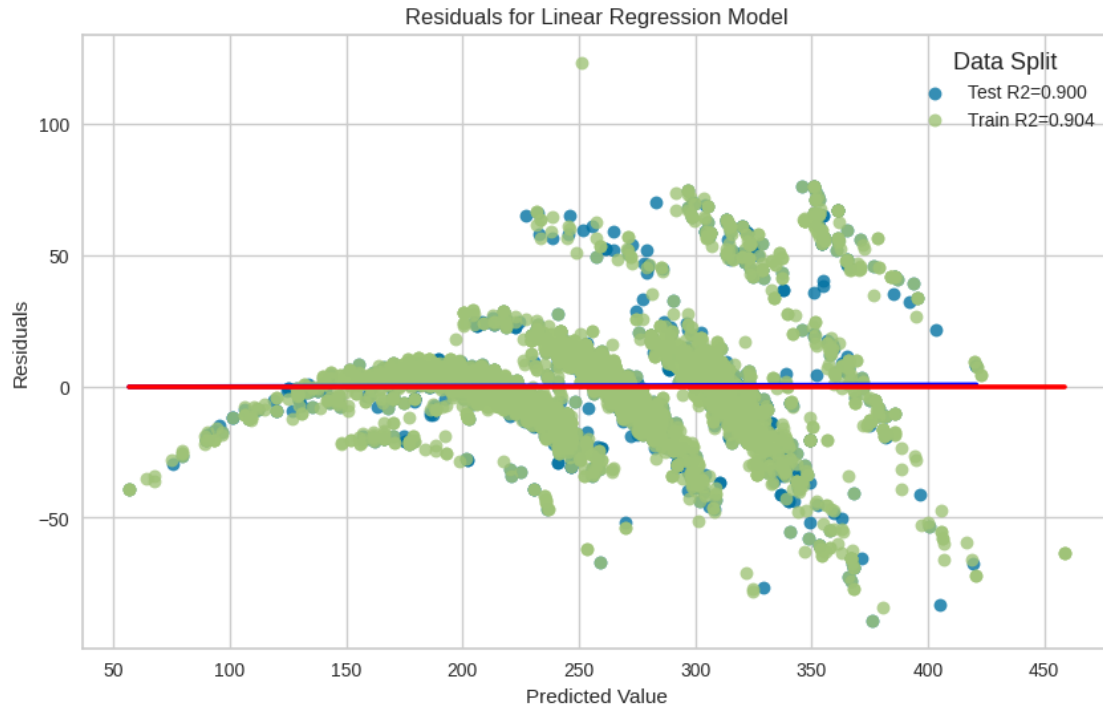
# Plot residuals for test data
sns.regplot(x=y_pred_test, y=(y_pred_test-y_test ), ci=None, label='Test R2=0.
↪900', scatter_kws={'s':50}, line_kws={'color':'blue'})

# Plot residuals for train data
sns.regplot(x=y_pred_train, y=(y_pred_train- y_train ), ci=None, label='Train_
↪R2=0.904', scatter_kws={'s':50}, line_kws={'color':'red'})

# Add title and labels
plt.title('Residuals for Linear Regression Model')
plt.xlabel('Predicted Value')
plt.ylabel('Residuals')

# Add legend with customization
plt.legend(loc='best', title='Data Split', title_fontsize='13', fontsize='10')

# Show the plot
plt.show()
```



- To determine how much the error deviates from the mean of the target label.
- According to the RMSE metric, our model has an average error rate of 7.3%.
- Prefer the RMSE metric because it penalizes poor predictions.
- the Multi linear regression model has good performance, with (R^2) values of 0.904 for the training set and 0.900 for the test set.
- Most residuals are close to zero, indicating accurate predictions.
- However, there are systematic errors as residuals increase with predicted values, suggesting the model may be biased for higher values.
- Additionally, the spread of residuals indicates the variance of errors is not constant. This suggests potential areas for model improvement.

the prediction error for a linear regression model with an (R^2) value of 0.900, indicating that 90% of the variance in the target variable is explained by the model. Overall, the model performs well with a high degree of accuracy.

Polynomial Features

```
[334]: from sklearn.preprocessing import PolynomialFeatures

def poly(d):

    train_rmse_errors = []
```

```

test_rmse_errors = []
number_of_features = []

for i in range(1, d):
    polynomial_converter = PolynomialFeatures(degree = i, include_bias=
↪False)
    poly_features = polynomial_converter.fit_transform(X)

    X_train, X_test, y_train, y_test = train_test_split(poly_features, y,
↪test_size=0.3, random_state=101)

    model = LinearRegression(fit_intercept=True)
    model.fit(X_train, y_train)

    train_pred = model.predict(X_train)
    test_pred = model.predict(X_test)

    train_RMSE = np.sqrt(mean_squared_error(y_train, train_pred))
    test_RMSE = np.sqrt(mean_squared_error(y_test, test_pred))

    train_rmse_errors.append(train_RMSE)
    test_rmse_errors.append(test_RMSE)

    number_of_features.append(poly_features.shape[1])

    return pd.DataFrame({"train_rmse_errors": train_rmse_errors,
↪"test_rmse_errors": test_rmse_errors,
                           "number of features": number_of_features},
↪index=range(1, d))

```

[336]: poly(10)

```

# The poly(10) function creates polynomial regression models of degrees 1 to
↪10,
# evaluates their training and test RMSE, and returns a DataFrame summarizing
↪these errors and
# the number of features used at each degree.

```

```

[336]:
train_rmse_errors  test_rmse_errors  number of features
1          18.314047          18.087420             4
2          15.788285          15.680230            14
3          14.314294          13.914205            34
4          12.916365          12.699119            69
5          12.093931          12.490365           125
6          11.139029          13.456785           209
7          19.939588         116.936156           329
8          12.676781         396.862449           494

```


This table shows RMSE and feature counts for polynomial regression models of different degrees.

Key insights:

- **Performance Improvement:** RMSE decreases up to the 6th degree, indicating better performance.
- **Optimal Degree:** The 6th degree has the lowest test RMSE (13.456785).
- **Overfitting:** The 7th degree shows overfitting with a high test RMSE (116.936156).
- **Feature Count:** The 4th degree model, with 69 features, achieves similar performance to the 5th degree model (125 features), making it more efficient and cost-effective.

The 4th degree is preferable for reducing computational cost while maintaining performance.

```
[343]: plt.plot(range(1,7), poly(7)["train_rmse_errors"], label = "TRAIN")
plt.plot(range(1,7), poly(7)["test_rmse_errors"], label = "TEST")
plt.xlabel("Polynomial Complex")
plt.ylabel("RMSE")
plt.legend()
plt.show()
```



Poly(degree=4)

[346]: # Selected degree=4

```
poly_converter = PolynomialFeatures(degree = 4, include_bias=False)
```

Model

[347]: # Poly linear model

```
poly_lin_reg = LinearRegression()
```

Train | Test Split

```
[348]: X_train, X_test, y_train, y_test = train_test_split(poly_converter.  
    ↪fit_transform(X), y,  
                                                    test_size = 0.2,  
    ↪random_state = 42)
```

Training the Model

[349]: poly_lin_reg.fit(X_train, y_train)

[349]: LinearRegression()

Predicting Test Data

```
[353]: y_pred_train = poly_lin_reg.predict(X_train)  
y_pred_test = poly_lin_reg.predict(X_test)
```

[354]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

```
R2Score_test=r2_score(y_test, y_pred_test)  
mae_test=mean_absolute_error(y_test, y_pred_test)  
mse_test=mean_squared_error(y_test, y_pred_test)  
rmse_test=np.sqrt(mean_squared_error(y_test, y_pred_test))
```

```
R2Score_train=r2_score(y_train, y_pred_train)  
mae_train=mean_absolute_error(y_train, y_pred_train)  
mse_train=mean_squared_error(y_train, y_pred_train)  
rmse_train=np.sqrt(mean_squared_error(y_train, y_pred_train))
```

```
print("performance Metrics of Multiple Linear Regression for Train and Test_  
    ↪data")  
print("\t\t Train_data   | Test_data")  
print("-----")  
print("R2 Score ",R2Score_train,"|",R2Score_test)
```



```

print("MAE      ",mae_train,"|",mae_test)
print("MSE      ",mse_train," |",mse_test)
print("RMSE     ",rmse_train,"|",rmse_test)

scores = {
    "MultiLinear_test": {"R2" : R2Score_test,
    "mae" : mae_test,
    "mse" : mse_test,
    "rmse" : rmse_test},

    "MultiLinear_train": {"R2" : R2Score_train,
    "mae" : mae_train,
    "mse" : mse_train,
    "rmse" : rmse_train}
}

PolyRegression_score=pd.DataFrame(scores)

PolyRegression_score

```

performance Metrics of Multiple Linear Regression for Train and Test data

	Train_data	Test_data
--	------------	-----------

R2 Score	0.953385177136985	0.9476458268934926
MAE	6.042713207467257	6.426362331953849
MSE	159.35515025782811	180.07884908237864
RMSE	12.623594981534701	13.419346075065604

```

[354]:      MultiLinear_test  MultiLinear_train
R2          0.947646         0.953385
mae          6.426362         6.042713
mse         180.078849        159.355150
rmse         13.419346        12.623595

```

Comparison of Simple Linear ,Multiple Linear & Poly Multiple Linear Regression

```

[357]: result = pd.concat([slr_score,Mlr_score, PolyRegression_score], axis=1)
result

```

```

[357]:      Simple_test  Simple_train  MultiLinear_test  MultiLinear_train \
R2          0.838150      0.844331         0.900108         0.903877
mae         14.233496      13.901692         11.514614         11.436689
mse         552.101088      533.438573        337.406607        330.829164
rmse         23.496831      23.096289         18.368631         18.188710

```

```

MultiLinear_test  MultiLinear_train

```

R2	0.947646	0.953385
mae	6.426362	6.042713
mse	180.078849	159.355150
rmse	13.419346	12.623595

0.6 Observation:

- The polynomial regression model (degree 4) performs better than the multiple linear and SimpleLinear regression model, achieving higher (R^2) values
- and lower MAE, MSE, and RMSE on both training and test sets, indicating better overall performance and accuracy.

[]: